

BOOKLY · CUSTOMER SUPPORT AGENT

The AI writes the answer. Code makes the decision.

Seyfullah Ural · September 2026

An AI wants to please you. A refund policy shouldn't.

A CHATBOT OPTIMISES FOR

**Making you happy
right now.**

A SUPPORT AGENT MUST OPTIMISE FOR

**Being right
about money.**

Both matter. Only one of them can be left to a language model.

Every answer takes the same three steps.

The middle one is the whole argument.

1

Claude picks a tool

It reads the message and chooses:
look up this order, check this refund,
search the help pages.

2

Python decides

45 days since delivery. The window is
30. The answer is no — computed by
code, not judged by the model.

3

Claude replies

It explains the outcome kindly. It is
not allowed to overturn step 2,
however the customer asks.

45 days old. The answer is no — every time.

A real exchange from the running agent.

I want a refund for BK-1002, ada.lovelace@example.com

I checked, and unfortunately BK-1002 isn't eligible for a refund — your copy of Piranesi was delivered on 23 July, 45 days ago, which falls outside the 30-day return window.

I'm sorry to give you that answer. If you'd like, I can pass this to a human colleague to take another look.

WHAT IT DID UNDERNEATH

```
check_return_eligibility → "eligible": false  "reason": "OUTSIDE_RETURN_WINDOW"
```

What I'd build next.

In order. The first one isn't close.

1 An eval set

The rules are covered by unit tests. The model's judgement isn't measured yet — so the next prompt change would be a guess.

2 Real session storage

Conversations live in memory today. A dropped connection shouldn't lose a customer.

3 Cost per conversation

Support is priced per contact. You can't argue for AI deflection without that number.