


BOOKLY · CUSTOMER SUPPORT AGENT

The AI writes the answer. Code makes the decision.

Seyfullah Ural · September 2026

How a question becomes an answer

Four moving parts. Only one of them guesses.



WHAT THE FLOW NEVER DOES

A customer turn enters on the left and leaves as a reply. Nothing in between ever asks the model what the rules are.

Memory is deliberately split. The model sees the transcript. Whether identity was verified is held by the application, in SQLite — the model cannot talk its way around a set.

Three decisions worth defending

Each one moves authority out of the model and into code.

01

**Eligibility is code,
not a prompt**

TRADED AWAY

Policy changes need a deploy, not a prompt edit.

WORTH IT BECAUSE

Auditable, unit-testable, and there is no prompt to jailbreak.

02

**The refund tool has
preconditions**

TRADED AWAY

More code than a single line of instruction.

WORTH IT BECAUSE

“Always confirm first” is a suggestion. This is enforced.

03

**Required arguments
make the dialogue**

TRADED AWAY

Less freedom in how the agent phrases its way there.

WORTH IT BECAUSE

Multi-turn collection with no dialogue script written.

The through-line: an AI wants to please you, and a refund policy shouldn't. Every decision above takes one more judgement away from the model.

Same refusal. Two languages.

Real replies from the running agent. English is the default — the customer just asked for German.

ENGLISH

"I want a refund for BK-1002"

BK-1002 isn't eligible for a refund — your copy of Piranesi was delivered on 23 July, 45 days ago, which falls outside the 30-day return window.

DEUTSCH

„Ich möchte eine Erstattung für BK-1002“

Leider kann ich für BK-1002 keine Erstattung veranlassen: Die Bestellung wurde vor 45 Tagen zugestellt und liegt damit außerhalb des 30-tägigen Rückgabefensters.

ONE TOOL RESULT — SHARED BY BOTH

```
check_return_eligibility → "eligible": false  "reason": "OUTSIDE_RETURN_WINDOW"
```

policy.py has never heard of German. It returns a code; turning that code into a sentence is a presentation job.

What I'd build next.

In order. The first one isn't close.

1 An eval set

The rules are covered by unit tests. The model's judgement isn't measured at all — so the next prompt change would be a guess.

2 Streaming replies

Answers arrive as a block after a pause. Streaming them word by word is the difference between a script and a product.

3 Cost per conversation

Support is priced per contact. You can't argue for AI deflection without that number.

Durable sessions and German were the previous two on this list. Both are built.